



COURSE CODE: BSDSESM1KU

BACHELOR IN SOFTWARE DEVELOPMENT

DevOps: ITU-MiniTwit

GROUP R — RHODODEVDRON

IT UNIVERSITY OF COPENHAGEN

Name	Email
Adrian Valdemar Borup	adbo@itu.dk
Albert Rise Nielsen	albn@itu.dk
Joachim Alexander Borup	aljb@itu.dk
Thomas Wolgast Rørbech	thwr@itu.dk

May 25, 2022

Table of Contents

1	System Perspective	2
1.1	System architecture	2
1.1.1	Backend	3
1.1.2	Frontend	3
1.1.3	Monitoring	3
1.1.4	Logging	3
1.1.5	Kubernetes and Azure	3
1.1.6	Supporting infrastructure	4
1.2	System design	4
1.2.1	Bridge and repository design patterns	4
1.2.2	Object-relational mapping	4
1.2.3	Separate frontend	5
1.2.4	Program modules	5
1.3	Software license compatibility	5
1.4	System dependencies	6
2	Process Perspective	6
2.1	Team interaction	6
2.2	Organisation of repositories	7
2.3	Git branching strategy	7
2.4	CI/CD chains	8
2.4.1	Building, testing, and static analysis	8
2.4.2	Deploying our service	8
2.4.3	Generating a changelog, and building the API \LaTeX project	8
2.5	Monitoring	9
2.6	Logging	9
2.7	Security assessment	10
2.8	Strategy for scaling and load balancing	11
	Appendices	i

1 System Perspective

1.1 System architecture

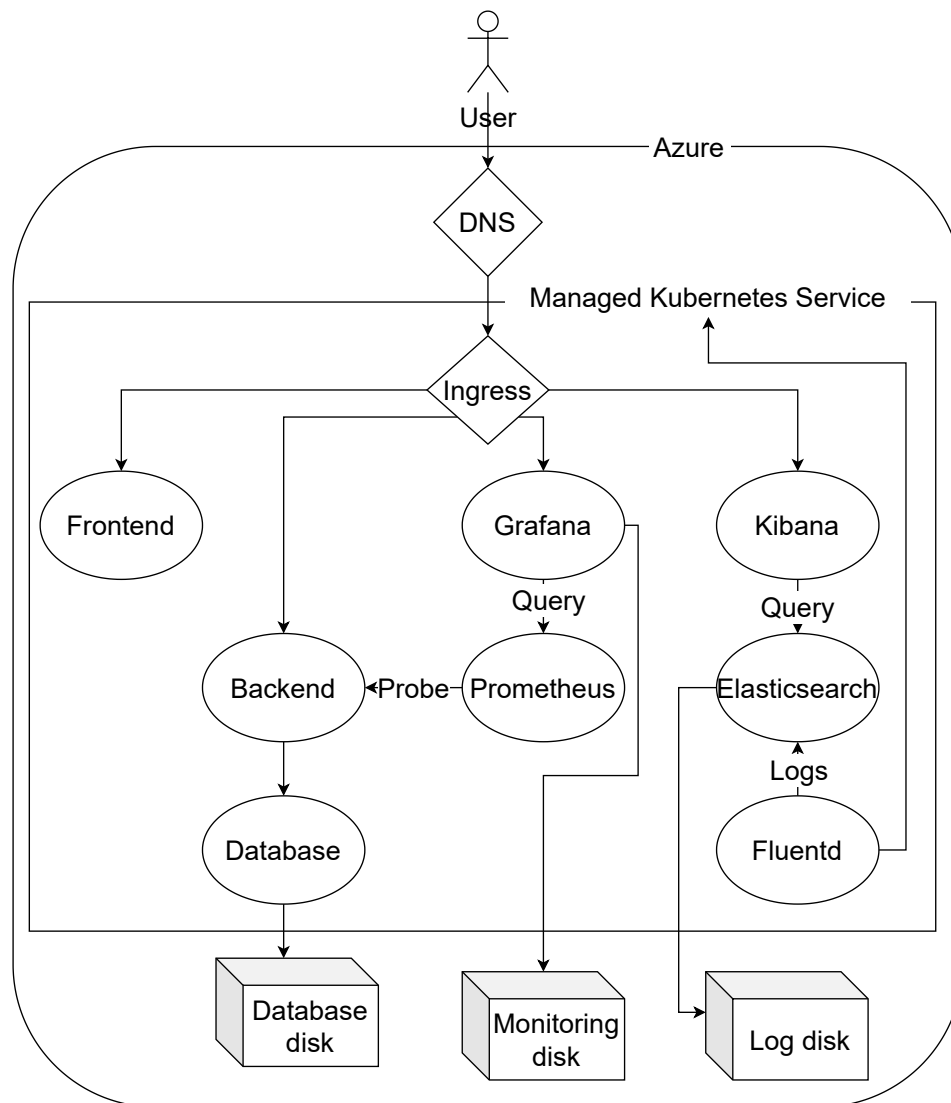


Figure 1: System architecture and its components

We've built our system twice, according to 2 specifications. The first was a, mostly, one to one rewrite of the original project into Go [tool:go]. While the second, and current iteration, was according to the API specification [spec:api]. This iteration consists of 4 subsystems and supporting infrastructure. The subsystems we'll go through below are: backend, frontend, monitoring and logging. The overview and interaction between these systems are modelled in figure 1.

1.1.1 Backend

The core system of the project is the backend. It is responsible for the actual computations and controlling of the database. It's built with Go [tool:go], runs in a Docker container and uses Microsoft SQL server [tool:microsoft-sql-server] as the database. We deep dive into this component in the "System Design" section - section 1.2.

1.1.2 Frontend

As the backend changed from server rendered website to a REST API, we decided to move the frontend into it's own subsystem. The frontend is built with Vue.js [tool:vue] and is written in TypeScript [tool:typescript]. It makes use of the public API that the backend exposes, and therefore provides a more user friendly access layer to our service. The subsystem is built in small components that can be reused when expanding the project. We had a focus on the visual elements being mostly identical to the original frontend, to ensure user familiarity.

1.1.3 Monitoring

To keep track of the health of our system, its performance, usage and other metrics, we use Prometheus [tool:prometheus] to collect metrics and Grafana [tool:grafana] to view said metrics. Our metrics include: CPU usage, memory usage, number of requests per endpoint, average request latency and amount of registered users. These metrics provide an overview over the system health and allows us to diagnose issues faster, by locating when something happened, at which endpoint and if it had an effect on system resources.

1.1.4 Logging

For more detailed diagnostics we have a logging subsystem. It's built with Fluentd [tool:fluentd], Elasticsearch [tool:elasticsearch], and Kibana [tool:kibana]. Fluentd reads the logs from the Kubernetes host, and parses them according to the configured log format. It then tags them and sends them to Elasticsearch, which indexes and stores them, such that they are easily searchable. Kibana provides a web interface to search, read and display the logs in a structured way.

Having this logging system allows us to diagnose what happened when something goes wrong, as we can easily find and read the info, warnings and errors that our system outputs.

1.1.5 Kubernetes and Azure

To deploy this system we use Kubernetes [tool:kubernetes] deployed on Azure [tool:azure] through AKS [tool:aks].

Using Kubernetes makes us host agnostic. A few times during the project we've had to move our host, which we've done without issue due to the emphasis on containerization and host abstraction that Kubernetes has. Kubernetes also allows us to easily scale our system, as we can easily deploy new nodes and load balance our services.

We've used Azure as our deployment target, as it was reliable, familiar and cheap. AKS is Azure's managed Kubernetes service.

1.1.6 Supporting infrastructure

To make this infrastructure function we have some supporting infrastructure. Such as our domain registrant where we bought our domain - one.com [**onecom**]. Our domain administrator that controls our domain and handles nameserver changes - DK Hostmaster [**dk-hostmaster**]. And our Azure hosted DNS zone [**tool:dns-zone**] that handles dns records. This informs systems to route traffic to the correct infrastructure.

1.2 System design

Our main system design goals were to make the system extensible and maintainable. That is, adding new features and replacing current implementations should be easy to do, and you should be able to trust that the system works once the change is made.

1.2.1 Bridge and repository design patterns

We've tried to achieve this by using the repository and bridge design patterns when possible, particularly to communicate with our database. The repository pattern is used to collect all our database queries in one place, and the bridge pattern is used to abstract away the implementation. Rather than talking directly to the database, our controllers communicate via an interface that defines which operations can be made. For an example of this, see the IUserRepository interface [**repo:user-repository-interface**].

Because we depend on the abstraction rather than the implementation, we can easily replace the database with a NoSQL database, create a wrapper that does extra operations on top, or similar—and the code depending on the interface won't have to change.

1.2.2 Object-relational mapping

In the same spirit as above, we've used GORM [**gorm**] for object-relational mapping. With this ORM, the database provider is abstracted away, making it possible to replace the database provider without changing the ORM code.

1.2.3 Separate frontend

We moved the frontend code to a separate project [**repo:frontend**]. This way, the Go server has only a single responsibility. This also serves as a way of using the bridge pattern to add another layer of abstraction, because the frontend is now only dependent on the API and not the Go implementation. This would also allow for different frontends to be implemented, like you obtain with MVC.

1.2.4 Program modules

In order to improve the single-responsibility aspect of our code, we've used modules to separate code into units of single responsibility as shown in Table 1.

Module	Responsibility
controllers	Managing the API routes and implementing the controllers
database	Communicate with the database
internal	Logging output destination and format — and error handling
models	Model definitions
monitoring	Managing Prometheus and monitoring
password	Password hashing and validation
test/controllers	Integration and unit tests for all controllers
main	Starting the application

Table 1: *The modules in our Go server and their responsibilities.*

1.3 Software license compatibility

To check if the software licenses of the libraries we've used are compatible with the MIT license in our repository, we have used lichen [**tool:lichen**] to automatically check the licenses of all used libraries. This process is documented in Appendix B. The results showed that all the libraries were compatible, except one because the tool couldn't find its license. However, checking its repository manually revealed a compatible MIT license.

1.4 System dependencies

Development stages	Tools and technologies	Description
Version control	Git, GitHub	Source code accessible for all developers, tracking issues, creating notes and tasks.
Pipeline CI/CD	CircleCi	Automate task such as; testing, building, deploying, releasing, generating changelog.
Containerization	Docker and DockerHub	Ensure same development environment and deployment to DockerHub.
Infrastructure	Terraform	IaC, Make it easier to swap host, automate tasks for releases.
Service Provider	Azure, own server	Hosting our app.
Service management	Kubernetes	Orchestrate containers, scaling up / down, load balancing.
Database	SQL Server, SQLite	To store user data.
Logging	Elasticsearch, Fluentd and Kibana	To gather logs into a centralized log and sort, filter, search the centralized log.
Monitoring	Grafana and Prometheus	Gather and organize metrics from our system, based on those create alerts and a visual representation.
Languages	Golang(insert cite), TypeScript(insert cite)	Backend and frontend.
Frameworks	gorm(insert cite), testify[tool:testify], gin-gonic, VueJS	ORM, Testing, HTTP Framework, Frontend SPA.
Code analysis	SonarCloud, lint, cyclomatic	Reduce code smell / Motivate clean code.
Security	Snyk, OWASP Zap, Metasploit WMAP	Notify about Security issues and updates to dependencies and scan for vulnerabilities.
Documentation	LaTeX, Markdown	Report, notes, changelog, instructions.

2 Process Perspective

2.1 Team interaction

The team consists of 4 developers, who meet once or twice a week to work on the project. At each meeting, the developers discuss the tasks at hand, create issues, divide the work, and then spends the rest of the time working on their tasks. There is no hierarchy - only 4 like-minded developers who makes decisions in unison.

New tasks are created as issues on GitHub, and then assigned to the developer who will be working on it. Maintaining a set of issues on GitHub lets us keep an overview of the planned tasks, and allows for any of the developers to easily add a new task at any time.

2.2 Organisation of repositories

The group uses GitHub to manage repositories. On GitHub, we have a dedicated organization that owns all of the project repositories. This way, we can keep related repositories together and reuse permission settings on GitHub. Currently, we have the following repositories:

1. Devops-2022-Group-R/itu-minitwit, which is the main project repository. Here, we have the backend web server which handles all business logic.
2. Devops-2022-Group-R/itu-minitwit-frontend, containing our frontend web-app.
3. Devops-2022-Group-R/flag-tool, which is a rewrite of the original flag tool given with the project template. The tool has been rewritten and moved to a separate repository.
4. Devops-2022-Group-R/bump-tool, which is a small tool to help with finding the next version bump based on a pull request's tag (major/minor/patch) and the project's current version.

The general philosophy has been to separate parts that can exist alone with a single responsibility. This way, all issues, pull requests, and releases are related to the topic of the repository. In a monolithic repository, you have to put in more effort to specify which part of the repository is relevant for a PR or issue. This also allows for easier and more simple CI/CD pipelines, and it becomes more obvious what a release changes. One example of this separation is the frontend web application, which is completely separated from the backend, with Kubernetes specs and CI/CD pipelines in its own repository.

However, we have not been as good at following this philosophy as we would like to. For example, our monitoring is a separate entity from our web server, but all the monitoring configuration is stored in the `itu-minitwit` repository [`repo:monitoring-config`] [`repo:monitoring-kubernetes`]. Instead, we should have moved this to a separate repository because it does not have anything to do with how the web server operates.

The same goes for the LaTeX files that this report consists of. We would have liked to have a separate repository for this in order to keep Git history clean and CI/CD more separate, but in this case it is a project requirement to have it in the main repository.

2.3 Git branching strategy

We have applied trunk-based development, meaning we branch out from the main branch, have a branch that lives shortly until the changes are implemented, and then merge it directly into the main branch. This keeps unrelated changes separate and makes PRs easier to review and merge. It's worth noting that we do not utilise release branches from trunk-based development, because we release continuously for each merge into the main branch.

2.4 CI/CD chains

We use CircleCI as a continuous integration and continuous delivery platform, and have configured a workflow [**workflow:circleci**] with jobs to build, test, and deploy our service, generate a changelog, as well as building our \LaTeX project. See Appendix A for a full CI/CD workflow diagram.

2.4.1 Building, testing, and static analysis

This part of the workflow quite simply installs the dependencies, and then builds and tests the project. Additionally it runs the following three static code analysis jobs:

- A Go linters runner: `golangci-lint`, which runs dozens of linters in parallel.
- A cyclomatic complexity calculator: `gocyclo`, to ensure that the code has low complexity. Otherwise, the job will exit with code 1 until we've refactored the functions in question.
- A developer security platform, Snyk, to find and fix security vulnerabilities in code, dependencies, and containers.

2.4.2 Deploying our service

CircleCI will also:

- Publish a GitHub release
- Push our changes to the four images in Docker Hub (the server, front-end, Prometheus, and Grafana)
- Deploy to their respective images in the managed Kubernetes service in the remote Azure server
- Apply infrastructure via Terraform

2.4.3 Generating a changelog, and building the \LaTeX project

We use CircleCI Github Changelog Generator [**tool:changelog-generator**] to generate a changelog based on our tags, issues, labels and pull requests on GitHub.

In order to include a build of the \LaTeX project in the CircleCI workflow, we've written a script to let CircleCI run a custom Docker image, build a PDF, and then commit it to the branch.

2.5 Monitoring

To monitor our systems, Prometheus is used to collect metrics and Grafana is used to visualize them. We chose to monitor the number of requests and request latency for each endpoint on our server, and we monitored the CPU and memory usage of it. We chose these because they give a quick overview of how the system is performing. The remaining information is stored in our logs.

For an example where we had an issue with our system, which was reflected in the monitoring, see Appendix C.




Figure 2: A sample of our Grafana dashboard.

2.6 Logging

To manage logging in our systems, we use the EFK stack (Elasticsearch, Fluentd, and Kibana). It works like this:

1. A pod in the cluster is logging to stdout
2. The stdout is sent to log files by Kubernetes
3. Fluentd picks up on logs in these files and sends them to Elasticsearch
4. Kibana is used to visualize the logs

We log using JSON, which allows us to select specific fields to display in Kibana, making the logs easy to read, filter, and sort by. Figure 3 shows how we display logs in an easy-to-read format.

Time ▾	kubernetes.container_name	method	route	status	duration
> Apr 5, 2022 @ 20:33:17.000	itu-minitwit-backend	POST	/fllws/Joan Exler?latest=2498757	204	53.309503ms
> Apr 5, 2022 @ 20:33:16.000	itu-minitwit-backend	POST	/fllws/Steffanie Kohut?latest=2498754	204	69.819242ms
> Apr 5, 2022 @ 20:33:16.000	itu-minitwit-backend	POST	/fllws/Jessie Mootispaw?latest=2498755	204	52.715469ms
> Apr 5, 2022 @ 20:33:16.000	itu-minitwit-backend	GET	/latest	200	1.210268ms
> Apr 5, 2022 @ 20:33:16.000	itu-minitwit-backend	POST	/fllws/Takako Mix?latest=2498756	204	69.464013ms
> Apr 5, 2022 @ 20:33:15.000	itu-minitwit-backend	POST	/fllws/Daphne Tarnowski?latest=2498751	204	58.167685ms
> Apr 5, 2022 @ 20:33:15.000	itu-minitwit-backend	POST	/fllws/Pearl Mccaskell?latest=2498752	204	47.46478ms
> Apr 5, 2022 @ 20:33:15.000	itu-minitwit-backend	POST	/fllws/Evelyn Lohmeier?latest=2498753	204	49.529896ms
> Apr 5, 2022 @ 20:33:14.000	itu-minitwit-backend	POST	/msgs/Major Zieglen?latest=2498750	204	79.164869ms
> Apr 5, 2022 @ 20:33:14.000	itu-minitwit-backend	GET	/latest	200	1.191267ms

Figure 3: Example of how we view logs in Kibana.

It's not shown on the figure, but we also have a column for error messages, which we've experienced makes it easy to diagnose problems in the system. To diagnose performance issues, the "duration" column has also been helpful. By the end of the project, the mentioned log output has been enough to diagnose bugs and get an overview of what's happening in the system.

2.7 Security assessment

We initially identified our assets by looking at which parts of our program might be vulnerable and of value to an intruder, we used the OWASP top 10 web application security risk [owasp-top-10] to identify sources of threats and from these we constructed risk scenarios to get a initial overview of the risks.

After having identified risks, we used this information during Risk analysis to create a risk matrix and discuss possible solutions for the different threats.

Some of the threats like SQL Injection and vulnerable or outdated dependencies were resolved by using an ORM [gorm] as middleware for the database to clean user input and static analysis tools integrated with our pipeline like Snyk [snyk] checking for vulnerable dependencies.

Risk Severity / Risk probability	Catastrophic	Critical	Marginal	Negligible	Insignificant
Certain					
Likely		DDOS			
Possible					
Unlikely			Access to Monitoring		
Rare	SQL Injection	Access Azure credentials	Access to CircleCi	Security flaw of vulnerable component	

Figure 4: Risk matrix

After risk analysis we used two vulnerability scanners OWASP ZAP [owasp-zap] and Metasploit [metasploit] WMAP [metasploit-wmap] to test our system for vulnerabilities tar-

getting our root endpoint [**minitwit-root-endpoint**]. We tried to add an extra middleware beside our CORS middleware to handle vulnerabilities flagged by the scanners, but we did not want to invest a lot of time and resources into it as we investigated the potential risks and realised they were just warnings and did not apply in our case. We dealt with one of our risk scenarios regarding monitoring, by ensuring changes to the dashboard were only accessible through code.

Read more in our security session notes [**repo:security-session-notes**].

2.8 Strategy for scaling and load balancing

We have deployed all of our components with Kubernetes deployments [**docs:kubernetes-deployment**], and configured them to run behind Kubernetes services [**docs:kubernetes-service**], this ensures component level load balancing. Using deployments also allows us to manually scale the number of instances of each component. While deployments do have the ability to autoscale, after some configuration, we have decided against it due to cost issues and because our monitoring indicated that the system was not getting enough load to justify scaling.


Deployments also support zero downtime updates. This is achieved by changing settings on the deployment, such as the image to be used, then Kubernetes creates a new pod, and then kills an old pod when the new is ready, repeating the process until all pods have the new configuration.

Appendices

Table of Contents

Appendix A	CI/CD workflow diagram	ii
Appendix B	Checking software license compatibility	iii
Appendix C	Monitoring example: High CPU usage	iv

A CI/CD workflow diagram



B Checking software license compatibility

To run the license scan tool, we used:

```
# Install
go install github.com/uw-labs/lichen@latest
# Run in root of Devops-2022-Group-R/itu-minitwit/
go build -o my-binary ./src
lichen my-binary
```

This reveals that the only disallowed library we use is Azure/azure-sdk-for-go/sdk/internal because it's missing a license. However, if you check the repository manually, you can find a MIT license in the root.

The output of the command is:

```
github.com/Azure/azure-sdk-for-go/sdk/azcore@v0.19.0: MIT (allowed)
github.com/Azure/azure-sdk-for-go/sdk/azidentity@v0.11.0: MIT (allowed)
github.com/Azure/azure-sdk-for-go/sdk/internal@v0.7.0: (not allowed - unresolvable license)
github.com/beorn7/perks@v1.0.1: MIT (allowed)
github.com/cespare/xxhash/v2@v2.1.2: MIT (allowed)
github.com/denisenkom/go-mssqldb@v0.12.0: BSD-3-Clause (allowed)
github.com/gin-contrib/sse@v0.1.0: MIT (allowed)
github.com/gin-gonic/gin@v1.7.7: MIT (allowed)
github.com/go-playground/locales@v0.14.0: MIT (allowed)
github.com/go-playground/universal-translator@v0.18.0: MIT (allowed)
github.com/go-playground/validator/v10@v10.10.0: MIT (allowed)
github.com/golang-sql/civil@v0.0.0-20190719163853-cb61b32ac6fe: Apache-2.0 (allowed)
github.com/golang-sql/sqlexp@v0.0.0-20170517235910-f1bb20e5a188: BSD-3-Clause (allowed)
github.com/golang/protobuf@v1.5.2: BSD-3-Clause (allowed)
github.com/jinzhu/inflection@v1.0.0: MIT (allowed)
github.com/jinzhu/now@v1.1.4: MIT (allowed)
github.com/leodido/go-urn@v1.2.1: MIT (allowed)
github.com/mattn/go-isatty@v0.0.14: MIT (allowed)
github.com/mattn/go-sqlite3@v1.14.11: MIT (allowed)
github.com/matttproud/golang_protobuf_extensions@v1.0.1: Apache-2.0 (allowed)
github.com/pkg/browser@v0.0.0-20180916011732-0a3d74bf9ce4: BSD-2-Clause (allowed)
github.com/prometheus/client_golang@v1.12.1: Apache-2.0 (allowed)
github.com/prometheus/client_model@v0.2.0: Apache-2.0 (allowed)
github.com/prometheus/common@v0.32.1: Apache-2.0 (allowed)
github.com/prometheus/procfs@v0.7.3: Apache-2.0 (allowed)
github.com/shirou/gopsutil@v3.21.11+incompatible: BSD-3-Clause (allowed)
github.com/sirupsen/logrus@v1.6.0: MIT (allowed)
github.com/tklauser/go-sysconf@v0.3.10: BSD-3-Clause (allowed)
github.com/tklauser/numcpus@v0.4.0: Apache-2.0 (allowed)
github.com/ugorji/go/codec@v1.2.6: MIT (allowed)
github.com/zsais/go-gin-prometheus@v0.1.0: MIT (allowed)
```

```

golang.org/x/crypto@v0.0.0-20220210151621-f4118a5b28e2: BSD-3-Clause (allowed)
golang.org/x/net@v0.0.0-20211112202133-69e39bad7dc2: BSD-3-Clause (allowed)
golang.org/x/sys@v0.0.0-20220209214540-3681064d5158: BSD-3-Clause (allowed)
golang.org/x/text@v0.3.7: BSD-3-Clause (allowed)
google.golang.org/protobuf@v1.27.1: BSD-3-Clause (allowed)
gopkg.in/yaml.v2@v2.4.0: Apache-2.0, MIT (allowed)
gorm.io/driver/sqlite@v1.2.6: MIT (allowed)
gorm.io/driver/sqlserver@v1.3.1: MIT (allowed)
gorm.io/gorm@v1.23.1: MIT (allowed)
2022/03/20 14:23:40 1 error occurred:
    * github.com/Azure/azure-sdk-for-go/sdk/internal@v0.7.0: not allowed - unresolvable
      license

```

C Monitoring example: High CPU usage

Figure 6 and Figure 7 show the status of our system in Grafana and Azure when we tried to implement logging. The initial implementation used up high amounts of our CPU resources, which is reflected in the response times and CPU recording metrics.

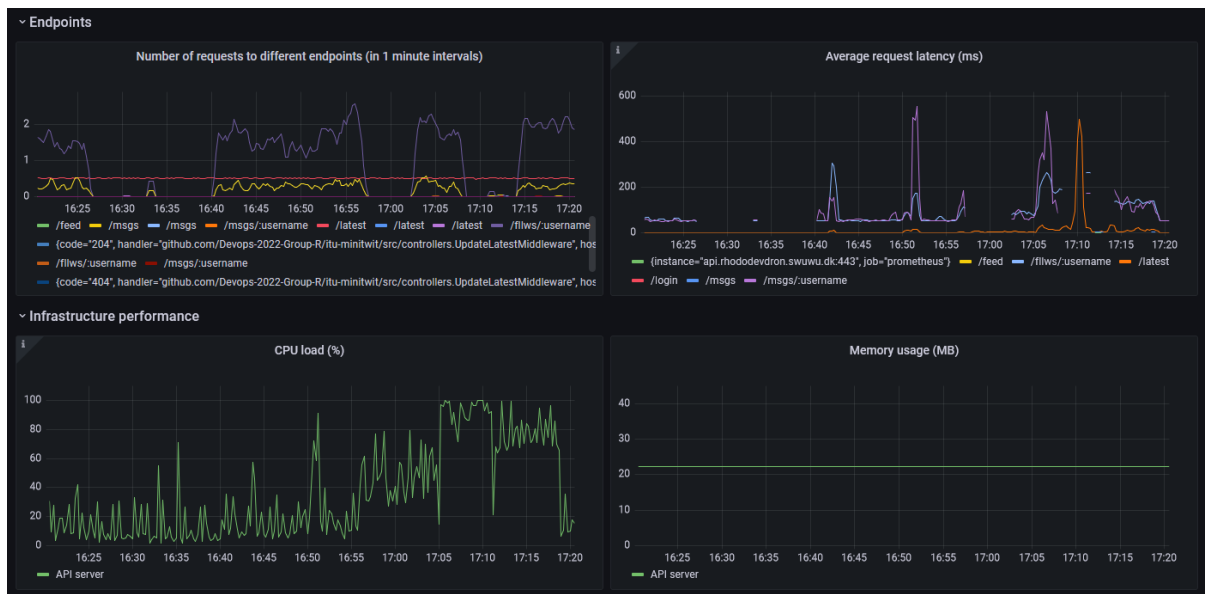


Figure 6: A snapshot of our Grafana monitoring dashboard.




Figure 7: A graph from Azure's monitoring showing the CPU usage of our cluster.